

# Choose a design pattern for your agentic AI system



## What is an Agent Design Pattern?

An **Agent Design Pattern** is a common architectural approach used to build agentic applications. It provides a distinct framework for:

1. **Organizing Components:** Structuring the agent's Model, Tools, and Memory.
2. **Integrating the Model:** Defining the role of the LLM (e.g., is it the sole planner or just a tool?).
3. **Orchestration:** Governing the flow and collaboration among single or multiple agents to accomplish a workflow.

### Why Agents Need Patterns

AI agents are best suited for:

- **Open-ended problems** requiring autonomous decision-making.
- **Complex multi-step workflow management.**
- **Knowledge-intensive tasks** that require real-time external data (grounding).

The patterns provide the structure needed to manage this complexity and autonomy safely and efficiently.



## The Design Process: Choosing a Pattern

Choosing the right design pattern is not a one-time decision but a process guided by your specific needs. It involves three high-level steps:

### 1. Define Your Requirements (Assess the Workload)

This initial step involves closely examining the characteristics of the task you want the agent to solve:

- **Task Complexity:** Is it a simple lookup (Level 1) or a multi-step planning and collaboration task (Level 3)?
- **Latency and Performance:** How quickly must the agent respond? (e.g., real-time customer service vs. nightly reporting).
- **Cost Budget:** More complex patterns involving multiple agent calls and large models cost more.
- **Need for Human Involvement (Safety):** Does the task involve high-stakes actions that require approval (e.g., financial transactions)?

## 2. Review the Common Agent Design Patterns

Once you know your requirements, you learn the capabilities of the available patterns (which include both **single-agent systems** and **multi-agent systems**).

## 3. Select a Pattern

You match the pattern's architecture and strengths to the workload characteristics defined in Step 1.

---

**Crucial Note:** This guide assumes you understand that agent architecture is fundamentally different from simpler AI applications like:

- **Direct Model Reasoning:** Calling the LLM once without tools or memory.
- **RAG (Retrieval-Augmented Generation):** Calling the LLM once with external documents (but usually without multi-step planning).

This foundational understanding ensures you choose an agent pattern only when the task truly requires the **autonomous decision-making** and **complex workflow management** that agents provide.

Here are your complete, structured notes on **Agentic AI Design Patterns**.

These are formatted specifically for you to copy and paste directly into a Google Doc for your portfolio or study guide.

---



# Agentic AI Design Patterns

**Definition:** Agent design patterns are architectural frameworks used to organize an AI system's components (Model, Tools, Memory) and orchestrate how they solve problems. Choosing the right pattern depends on the trade-off between

**Autonomy** (flexibility) and **Control** (reliability).

---

## 1. Foundational Architectures

### Single-Agent System

The fundamental building block. A single AI model uses a defined set of tools and a comprehensive system prompt to handle a request from start to finish.

- **How it works:** The user sends a prompt → The Agent reasons → Calls Tools → Returns Answer.
- **Best For:** Prototypes, MVPs, and tasks requiring simple tool use (e.g., "Search for the weather and convert to Celsius").
- **Pros:** Simple to build, low cost, easy to debug.
- **Cons:** Fails as complexity grows; context window gets overcrowded; single point of failure.

### Multi-Agent System (General Concept)

A system that orchestrates multiple **specialized agents** to solve a complex problem.

- **Core Principle: Decomposition.** Breaking a massive goal into smaller sub-tasks (e.g., Researcher, Writer, Coder).
- **Why Use It:** When a single prompt is too complex for one context window, or when different sub-tasks require different tools/personas.
- **Key Component: Context Engineering** (managing what information is passed between agents).

## 2. Deterministic Workflows (Predictable)

*Use when the workflow path is known in advance and does not change.*

### Sequential Pattern (The Assembly Line)

Executes agents in a fixed, linear order.

- **Structure:** Input → [Agent A] → Output A/Input B → [Agent B] → Output.
- **Best For:** Highly structured pipelines (e.g., Data Extraction → Data Cleaning → Database Insertion).
- **Trade-off:** Low latency and cost, but very rigid. If one step fails or is slow, the whole chain stops.

### Parallel Pattern (Concurrent Processing)

Executes multiple sub-agents at the same time, then synthesizes the results.

- **Structure:** Input is sent simultaneously to Agent A, Agent B, and Agent C. Their outputs are aggregated by a final step.

- **Best For:** Tasks needing diverse perspectives or speed (e.g., Analyzing a stock symbol with Sentiment Agent, Technical Agent, and Fundamental Agent simultaneously).
- **Trade-off:** Reduces latency (speed), but increases cost (multiple calls at once) and complexity in merging conflicting results.

### Loop Pattern (The Monitor)

Repeatedly executes a task until a specific condition is met.

- **Structure:** Run Agent → Check Exit Condition. If "No", run again. If "Yes", stop.
- **Best For:** Monitoring tasks or retries (e.g., "Check server status every 5 minutes until it returns 200 OK").
- **Trade-off:** Risk of infinite loops (high cost) if the exit condition is never met.

## 3. Iterative & Quality Workflows (The "Quality" Loop)

*Use when the first draft is likely imperfect and needs refinement.*

### ReAct Pattern (Reason + Act)

The standard operating system for modern agents. It forces the model to "think" before it "acts."

- **Structure:**
  1. **Thought:** Model analyzes the request.
  2. **Action:** Model chooses a tool.
  3. **Observation:** Model reads the tool output.
  4. **Repeat:** Updates thought based on observation until done.
- **Best For:** Dynamic tasks requiring continuous planning (e.g., "Find the phone number of the CEO of the company that makes the iPhone").
- **Visual:**

### Review and Critique Pattern (Generator & Critic)

A specific quality assurance workflow using two distinct agents.

- **Structure:**
  1. **Generator Agent:** Creates the content (e.g., Code).

2. **Critic Agent:** Evaluates it against hard criteria (e.g., "Does it have security bugs?").
  3. **Feedback:** If rejected, sends feedback back to Generator.
- **Best For:** High-stakes content generation (Code generation, Legal drafting) where accuracy is paramount.
  - **Trade-off:** Higher cost (paying for two agents per turn).

### Iterative Refinement Pattern (The Polisher)

The general strategy of progressively improving an output through cycles.

- **Structure:** Generate Draft → Evaluate → Enhance Prompt/Context → Regenerate Draft.
  - **Advanced Feature:** Often uses a **Prompt Enhancer** subagent to rewrite the instructions for the next loop based on previous failures.
  - **Best For:** Complex creative or reasoning tasks where the "perfect" answer requires polishing (e.g., writing a novel chapter).
  - **Visual:**
- 

## 4. Dynamic Orchestration (The "Manager" Models)

*Use when the path is unknown and requires AI intelligence to navigate.*

### Coordinator Pattern (The Router)

Uses a central "Manager" agent to dynamically route tasks.

- **Structure:** User → **Coordinator Agent** (Decides "Who handles this?") → Routes to **Specialist Subagent**.
- **Best For:** Adaptive workflows like Customer Support (Classifying "Refund" vs. "Tech Support" vs. "Sales" and routing accordingly).
- **Trade-off:** Flexible and smart, but adds latency (the Manager must "think" before anyone works).
- **Visual:**

### Hierarchical Task Decomposition (The Organization Chart)

A multi-level hierarchy of agents handling massive planning.

- **Structure:** Root Agent (CEO) breaks goal into chunks → Middle Agents (Managers) break chunks into tasks → Leaf Agents (Workers) execute tasks.
- **Best For:** Extremely ambiguous, large-scale goals (e.g., "Plan a 3-day conference," "Write a software application from scratch").
- **Trade-off:** Most powerful, but highest latency and hardest to debug.
- **Visual:**

### Swarm Pattern (The Roundtable)

Collaborative, all-to-all communication without a central boss.

- **Structure:** Agents act as peers. Agent A can talk to B, who talks to C, who talks back to A. They debate and hand off tasks dynamically.
- **Best For:** Brainstorming, creative problem solving, and tasks requiring diverse expert viewpoints without a rigid process.
- **Trade-off:** High creativity, but high risk of "looping" (arguing forever) and very expensive.

## 5. Special Patterns (Safety)

### Human-in-the-Loop (HITL)

Integrates a mandatory checkpoint for human approval.

- **Structure:** Agent works → Reaches Checkpoint → Pauses State → Sends Notification (Email/UI) → Waits for Human Click → Resumes.
- **Best For:** Sensitive actions (Refunds > \$100, deploying code to production, deleting database records).
- **Visual:**



### Quick Decision Matrix (Cheatsheet)

If your workflow is...	Use this Pattern
------------------------	------------------

<b>Predictable &amp; Linear</b>	<b>Sequential</b>
<b>Needs Speed / Independent Tasks</b>	<b>Parallel</b>
<b>Requires Monitoring/Retries</b>	<b>Loop</b>
<b>Requires Quality/Safety</b>	<b>Review &amp; Critique or HITL</b>
<b>Needs Dynamic Routing</b>	<b>Coordinator</b>
<b>Huge, Ambiguous Project</b>	<b>Hierarchical Decomposition</b>
<b>Creative Brainstorming</b>	<b>Swarm</b>